

Monika Kusumanchi

AI Engineer | +91 6281074516 | kusumonika033@gmail.com

Summary

AI Engineer specializing in LLM systems, conversational AI, and scalable ML infrastructure with 3+ years building production-grade AI copilots, RAG pipelines, and voice agents. Reduced inference costs by 60% and delivered low-latency AI systems (<10s response) at scale on AWS.

Core Skills

LLMs, RAG, LangChain, LangGraph, OpenAI, Anthropic, Gemini • Python, TypeScript, FastAPI, Node.js • AWS (Bedrock, Lambda, ECS), Docker, Kubernetes • Vector DBs (Pinecone, Weaviate, pgvector), Neo4j • Voice AI (Twilio, Deepgram, ElevenLabs) • PostgreSQL, Redis, Kafka

Experience

AI Engineer — Recent Role

Built production AI copilots and RAG pipelines; cut inference costs 60% via prompt + model optimization; shipped voice agents with sub-10s response latency.

ML Engineer

Designed scalable ML infra on AWS; deployed end-to-end LLM features; owned eval + observability stack.

Data / AI Roles (earlier)

Analytics, NLP pipelines, model fine-tuning, internal tooling.

Achievements

• 3+ hackathon wins • Winner of NLP competition (200+ participants) • Best Performer award for AI call analytics optimization

Certifications

Agentic AI (DeepLearning.AI) • Microsoft PowerBI Data Analyst • Introduction to Data Engineering • Neo4j Certified Professional • Claude Code Certification